

YData: An Introduction to Data Science

Lecture 08: Histograms

Jessi Cisewski-Kehe and John Lafferty
Statistics & Data Science, Yale University
Spring 2019

Credit: data8.org



Announcements

Types of Data

Two Important Types

All values in a column should be both the same type and be comparable to each other in some way

- **Numerical** – Each value is from a numerical scale
 - Ordered, because they are numbers
 - Differences, averages, etc are meaningful
- **Categorical** – Each value is from a fixed inventory
 - May or may not have an ordering

Terminology

- **Individuals**: those whose features are recorded
- **Variable**: a feature, an attribute
- A variable has different **values**
- Values can be **numerical** or **categorical**, and of many sub-types within these
- Each **individual has exactly one value** of the variable
- **Distribution**: For each different value of the variable, the frequency of individuals that have that value

Categorical Distributions

- Bar charts are commonly used to visualize categorical distributions
- One axis is categorical, one numerical

(DEMO)

Displaying a Categorical Distribution


- The distribution of a variable (a column, e.g. Studios) describes the frequencies of its different values
- The **group** method counts the number of rows for each value in the column (e.g. the number of top movies released by each studio)
- Bar charts can display the distribution of a categorical variable (e.g. studios):
 - One bar for each category
 - Length of bar is the count of individuals in that category

Binning a Numerical Variable

Binning Numerical Values


Binning is counting the number of numerical values that lie within ranges, called bins.

- Bins are defined by their lower bounds (inclusive)
- The upper bound is the lower bound of the next bin



Area Principle

What Is Wrong With This Picture?



From [Gizmodo](http://gizmodo.com/new-ipad-battery-size-is-huge/), this shows battery size in the new iPad versus that of the iPad 2. The battery in the former is 70 percent bigger than that of the latter. Something's not right here.

Image credit: <http://flowingdata.com/2012/03/16/new-ipad-battery-size-is-huge/>

Area Principle

Areas should be proportional to the values they represent.

For example

- If you represent 20% of a population by



- Then 40% can be represented by:



- But not by:



Drawing Histograms

Histogram

- Chart that displays the distribution of a numerical variable
- Uses bins; there is one bar corresponding to each bin
- Uses the area principle:
 - The area of each bar is the percent of individuals in the corresponding bin

(DEMO)

Density

Histogram Axes

- By default, `hist` uses a scale (`normed=True`) that ensures the area of the chart sums to 100%
- The area of each bar is a percentage of the whole
- The horizontal axis is a number line (e.g., years), and the bins sizes don't have to be equal to each other
- The vertical axis is a rate (e.g., percent per year)

(DEMO)

How to Calculate Height

The [40, 65) bin contains 52 out of 200 movies

- “52 out of 200” is 26%
- The bin is $65 - 40 = 25$ years wide

$$\begin{aligned}\text{Height of bar} &= \frac{26 \text{ percent}}{25 \text{ years}} \\ &= 1.04 \text{ percent per year}\end{aligned}$$

$$\text{Height of bar} = \frac{\% \text{ in bin}}{\text{width of bin}}$$

- The height measures the percent of data in the bin **relative to the amount of space in the bin.**
- Height measures crowdedness, or **density**.
- Units: percent per unit on the horizontal axis

Area Measures Percent

Area of bar = % in bin = Height \times width of bin

- “How many individuals in the bin?” Use **area**.
- “How crowded is the bin?” Use **height**.

Bar Chart or Histogram?

To display a distribution:

Bar Chart

Distribution of categorical variable

Bars have arbitrary (but equal) widths and spacings

height (or length) of bars proportional to the percent of individuals

Histogram

Distribution of numerical variable

Horizontal axis is numerical: to scale, no gaps, bins can be unequal

Area of bars proportional to the percent of individuals; height measures density

Discussion Questions

What is the height of each bar in this histogram?

```
incomes.hist(1, bins=[0,15,25,85])
```

What are the vertical axis units?

Name	2016 Income (millions)
Jennifer Lawrence	61.7
Scarlett Johansson	57.5
Angelina Jolie	40
Jennifer Aniston	24.75
Anne Hathaway	24
Melissa McCarthy	24
Bingbing Fan	20
Sandra Bullock	20
Cara Delevingne	15
Reese Witherspoon	15
Amy Adams	15
Kristen Stewart	12
Amanda Seyfried	10.5
Tina Fey	10.5
Julia Roberts	10
Emma Stone	10
Natalie Portman	8.5
Margot Robbie	8
Meryl Streep	6
Mila Kunis	4.5

Vertical axis units: Percent per million

```
incomes.hist(1, bins=[0,15,25,85])
```

$[0, 15)$: $(45\%)/(15 \text{ million})$
= 3 % per million

$[15, 25)$: $(40\%)/(10 \text{ million})$
= 4 % per million

$[25, 85)$: $(15\%)/(60 \text{ million})$
= 0.25 % per million

Name	2016 Income (millions)
Jennifer Lawrence	61.7
Scarlett Johansson	57.5
Angelina Jolie	40
Jennifer Aniston	24.75
Anne Hathaway	24
Melissa McCarthy	24
Bingbing Fan	20
Sandra Bullock	20
Cara Delevingne	15
Reese Witherspoon	15
Amy Adams	15
Kristen Stewart	12
Amanda Seyfried	10.5
Tina Fey	10.5
Julia Roberts	10
Emma Stone	10
Natalie Portman	8.5
Margot Robbie	8
Meryl Streep	6
Mila Kunis	4.5